

Descriptive Analytics

Melissa Paniagua

10/15/2020

Assignment 1 | Module 4

The purpose of this assignment is to develop descriptive statistics analysis utilizing data from an online retail company.

The data contains the following attributes:

- InvoiceNo: Invoice number has a 6-digit integral number uniquely assigned to each transaction. If this code starts with letter 'c', it indicates a cancellation.
- StockCode: Product code. It is a 5-digit integral number uniquely assigned to each distinct product.
- Description: Product name.
- Quantity: The quantities of each product per transaction.
- InvoiceDate: It shows the day and time when each transaction was generated.
- UnitPrice: Product price per unit in sterling.
- CustomerID: It is a 5-digit integral number uniquely assigned to each customer.
- Country: The name of the country where each customer resides.

```
# Load the libraries needed for this assignment
```

```
library(readr)
```

```
library(dplyr)
```

```
library(tidyverse)
```

```
#Load the dataset
```

```
mydf <- read.csv("Online_Retail.csv")
```

```
#See the first 6 columns of the dataset
```

```
head(mydf)
```

	InvoiceNo	StockCode	Description	Quantity
1	536365	85123A	WHITE HANGING HEART T-LIGHT HOLDER	6
2	536365	71053	WHITE METAL LANTERN	6
3	536365	84406B	CREAM CUPID HEARTS COAT HANGER	8
4	536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6
5	536365	84029E	RED WOOLLY HOTTIE WHITE HEART.	6
6	536365	22752	SET 7 BABUSHKA NESTING BOXES	2
	InvoiceDate	UnitPrice	CustomerID	Country

1	12/1/2010 8:26	2.55	17850 United Kingdom
2	12/1/2010 8:26	3.39	17850 United Kingdom
3	12/1/2010 8:26	2.75	17850 United Kingdom
4	12/1/2010 8:26	3.39	17850 United Kingdom
5	12/1/2010 8:26	3.39	17850 United Kingdom
6	12/1/2010 8:26	7.65	17850 United Kingdom

Data Exploration

```
#To get descriptive statistics
summary(mydf)
```

InvoiceNo	StockCode	Description
573585 : 1114	85123A : 2313	WHITE HANGING HEART T-LIGHT HOLDER: 2369
581219 : 749	22423 : 2203	REGENCY CAKESTAND 3 TIER : 2200
581492 : 731	85099B : 2159	JUMBO BAG RED RETROSPOT : 2159
580729 : 721	47566 : 1727	PARTY BUNTING : 1727
558475 : 705	20725 : 1639	LUNCH BAG RED RETROSPOT : 1638
579777 : 687	84879 : 1502	ASSORTED COLOUR BIRD ORNAMENT : 1501
(Other):537202	(Other):530366	(Other) :530315

Quantity	InvoiceDate	UnitPrice
Min. : -80995.00	10/31/2011 14:41: 1114	Min. : -11062.06
1st Qu.: 1.00	12/8/2011 9:28 : 749	1st Qu.: 1.25
Median : 3.00	12/9/2011 10:03 : 731	Median : 2.08
Mean : 9.55	12/5/2011 17:24 : 721	Mean : 4.61
3rd Qu.: 10.00	6/29/2011 15:58 : 705	3rd Qu.: 4.13
Max. : 80995.00	11/30/2011 15:13: 687	Max. : 38970.00
	(Other) :537202	

CustomerID	Country
Min. :12346	United Kingdom:495478
1st Qu.:13953	Germany : 9495
Median :15152	France : 8557
Mean :15288	EIRE : 8196
3rd Qu.:16791	Spain : 2533
Max. :18287	Netherlands : 2371
NA's :135080	(Other) : 15279

Here we can see that CustomerID variable has 135,080 missing values.

```
# To see the number of data points
nrow(mydf)
```

```
[1] 541909
```

```
# Number of transactions for each country
head(table(mydf$Country))
```

Australia	Austria	Bahrain	Belgium	Brazil	Canada
1259	401	19	2069	32	151

Questions:

1. Show the breakdown of the number of transactions by countries i.e. how many transactions are in the dataset for each country (consider all records including cancelled transactions). Show this in total number and also in percentage. Show only countries accounting for more than 1% of the total transactions.

```
#Total Number of transaccions for each country showing 1% of the total transactions
mypct <- mydf %>% group_by(Country) %>%
  summarise(Total_Trans = n(), Total_Perc = sum(n()/length(mydf$Country)*100)) %>%
  filter(Total_Perc > 1)

# See the dataframe
as.data.frame(mypct)
```

	Country	Total_Trans	Total_Perc
1	EIRE	8196	1.512431
2	France	8557	1.579047
3	Germany	9495	1.752139
4	United Kingdom	495478	91.431956

2. Create a new variable 'TransactionValue' that is the product of the existing 'Quantity' and 'UnitPrice' variables. Add this variable to the dataframe.

```
# Include the new variable
mydf <- mydf %>% mutate(TransactionValue = Quantity * UnitPrice)

# See the first 6 rows and last 6 columns of the dataframe
mydf[1:6, 4:9]
```

	Quantity	InvoiceDate	UnitPrice	CustomerID	Country	TransactionValue
1	6	12/1/2010 8:26	2.55	17850	United Kingdom	15.30
2	6	12/1/2010 8:26	3.39	17850	United Kingdom	20.34
3	8	12/1/2010 8:26	2.75	17850	United Kingdom	22.00
4	6	12/1/2010 8:26	3.39	17850	United Kingdom	20.34
5	6	12/1/2010 8:26	3.39	17850	United Kingdom	20.34
6	2	12/1/2010 8:26	7.65	17850	United Kingdom	15.30

3. Using the newly created variable, TransactionValue, show the breakdown of transaction values by countries i.e. how much money in total has been spent each country. Show this in total sum of transaction values. Show only countries with total transaction exceeding 130,000 British Pound.

```
# Show the countries and total transactions greater than 130,000 British Pound
greater_trans <- mydf %>% select(Country, TransactionValue) %>%
  group_by(Country) %>%
  summarise(Transactions = sum(TransactionValue)) %>%
  filter(Transactions > 130000)

# See the values
as.data.frame(greater_trans)
```

	Country	Transactions
1	Australia	137077.3
2	EIRE	263276.8
3	France	197403.9
4	Germany	221698.2
5	Netherlands	284661.5
6	United Kingdom	8187806.4

4. Convert 'InvoiceDate' from categorical into date variable:

```
#First let's convert 'InvoiceDate' into a POSIXlt object
temp <- strptime(mydf$InvoiceDate,format='%m/%d/%Y %H:%M',tz='GMT')

# See the dataframe
head (temp)
```

```
[1] "2010-12-01 08:26:00 GMT" "2010-12-01 08:26:00 GMT"
[3] "2010-12-01 08:26:00 GMT" "2010-12-01 08:26:00 GMT"
[5] "2010-12-01 08:26:00 GMT" "2010-12-01 08:26:00 GMT"
```

```
#Now, separate date, day of the week, and hour components dataframe
mydf$New_Invoice_Date <- as.Date(temp)

# Know the difference between the two dates in terms of the number days.
mydf$New_Invoice_Date[20000] - mydf$New_Invoice_Date[10]
```

Time difference of 8 days

```
# Define a new variable with the day name
mydf$Invoice_Day_Week = weekdays(mydf$New_Invoice_Date)

# Convert the hour into a normal numerical value
mydf$New_Invoice_Hour = as.numeric(format(temp, "%H"))

#Finally, define the month as a separate numeric variable.
mydf$New_Invoice_Month = as.numeric(format(temp, "%m"))

# To see the dataframe with the new columns
mydf[1:6, 10:13]
```

	New_Invoice_Date	Invoice_Day_Week	New_Invoice_Hour	New_Invoice_Month
1	2010-12-01	Wednesday	8	12
2	2010-12-01	Wednesday	8	12
3	2010-12-01	Wednesday	8	12
4	2010-12-01	Wednesday	8	12
5	2010-12-01	Wednesday	8	12
6	2010-12-01	Wednesday	8	12

- a. Show the percentage of transactions (by numbers) by days of the week.

```
# To get the total number of days transactions and its percentage
perc_day <- mydf %>% group_by(Invoice_Day_Week) %>%
  summarise(Num_Trans = n(), Percent = sum(n())/length(mydf$Invoice_Day_Week)*100))

# Show the dataframe
as.data.frame(perc_day)
```

	Invoice_Day_Week	Num_Trans	Percent
1	Friday	82193	15.16731
2	Monday	95111	17.55110
3	Sunday	64375	11.87930
4	Thursday	103857	19.16503
5	Tuesday	101808	18.78692
6	Wednesday	94565	17.45035

- b) Show the percentage of transactions (by transaction volume) by days of the week.

```
# To get the total volume of transactions and its percentage per week days
total_perc_day <- mydf %>% group_by(Invoice_Day_Week) %>%
  summarise(Total_Trans = sum(TransactionValue)) %>%
  mutate(Percent = Total_Trans/sum(Total_Trans)*100)

# Show the dataframe
as.data.frame(total_perc_day)
```

	Invoice_Day_Week	Total_Trans	Percent
1	Friday	1540610.8	15.804787
2	Monday	1588609.4	16.297194
3	Sunday	805678.9	8.265282
4	Thursday	2112519.0	21.671867
5	Tuesday	1966182.8	20.170636
6	Wednesday	1734147.0	17.790232

- c) Show the percentage of transactions (by transaction volume) by month of the year.

```
# To get the total volume of transactions by month
perc_month <- mydf %>% group_by(New_Invoice_Month) %>%
  summarise(Total_Trans = sum(TransactionValue)) %>%
  mutate(Percent = Total_Trans/sum(Total_Trans)*100)

# Show the dataframe
as.data.frame(perc_month)
```

	New_Invoice_Month	Total_Trans	Percent
1	1	560000.3	5.744919
2	2	498062.6	5.109515
3	3	683267.1	7.009487
4	4	493207.1	5.059703
5	5	723333.5	7.420519

6	6	691123.1	7.090080
7	7	681300.1	6.989308
8	8	682680.5	7.003469
9	9	1019687.6	10.460751
10	10	1070704.7	10.984123
11	11	1461756.2	14.995836
12	12	1182625.0	12.132290

- d) What was the date with the highest number of transactions from Australia?

```
# To select the highest number of transactions per country
highest_num1 <- mydf %>%
  filter(mydf$Country == "Australia") %>%
  group_by(New_Invoice_Date) %>%
  summarise(Num_TransactionValue = n()) %>%
  top_n(1, Num_TransactionValue)


# Show the dataframe
as.data.frame(highest_num1)
```

	New_Invoice_Date	Num_TransactionValue
1	2011-06-15	139

- e) The company needs to shut down the website for two consecutive hours for maintenance. What would be the hour of the day to start this so that the distribution is at minimum for the customers? The responsible IT team is available from 7:00 to 20:00 every day.

```
# Histogram plot to show the less busy time of the day to develop the maintenance.
hist(mydf$New_Invoice_Hour,
     main="Histogram for Hours",
     xlab="Hours",
     border="darkblue",
     col="lightblue",
     xlim=c(7,20),
     ylab="Frequency of Hours",
     las=1,
     breaks=12)

# Bar plot to identify the best days to implement the maintenance
barplot(table(mydf$Invoice_Day_Week),
        main= "Bar Chart of Day",
        col= c("darkblue"))
```



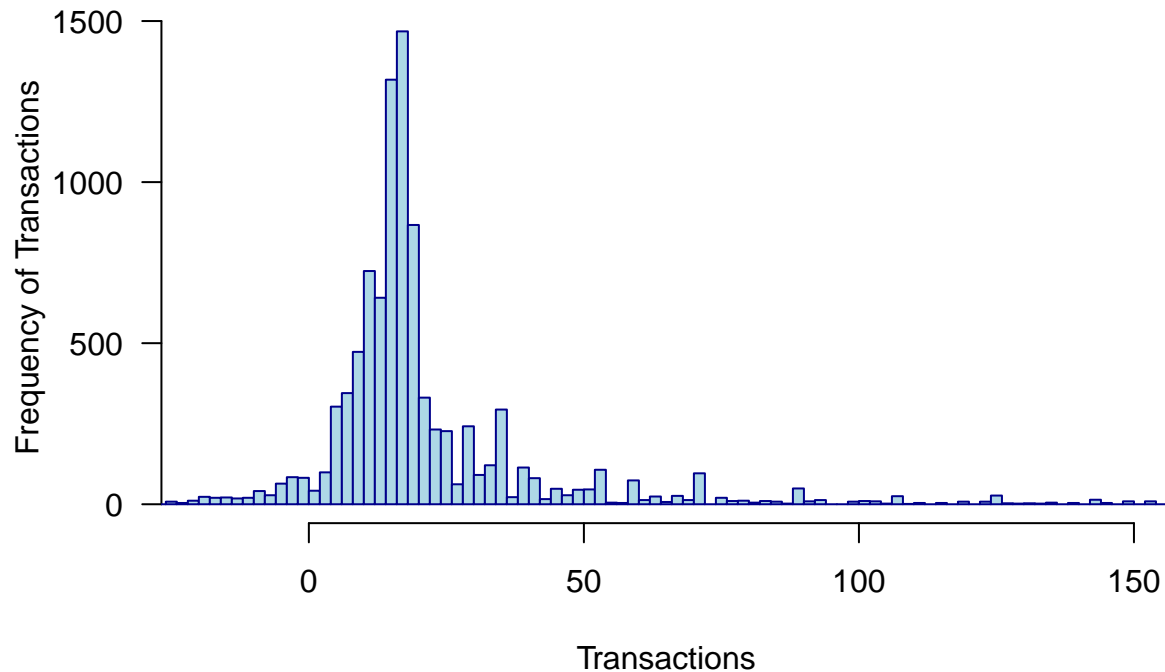
Here we can see on this plot, the best hours to do the maintenance of the company's website are between 18:00 and 20:00 pm. Additionally, Sundays night would be a good day to do the maintenance of the website.

5. Plot the histogram of transaction values from Germany. Use the `hist()` function to plot.

```
# Transaction values from Germany
germany <- select(mydf, TransactionValue, Country) %>%
  filter(mydf$Country == "Germany")

# Histogram
hist(germany$TransactionValue,
     main="Histogram of Germany's Transaction Values",
     xlab="Transactions",
     border="darkblue",
     col="lightblue",
     xlim=c(-20, 150),
     ylab="Frequency of Transactions",
     las=1,
     breaks=600)
```

Histogram of Germany's Transaction Values



6. Which customer had the highest number of transactions? Which customer is most valuable (i.e. highest total sum of transactions)?

Part a)

```
# Most Valuable Customer with highest number of transactions
valuable_customer <- mydf %>% na.omit() %>%
  group_by(CustomerID) %>%
  summarise(Highest_Num = n()) %>%
  top_n(1, Highest_Num)

# Show the dataframe
as.data.frame(valuable_customer)
```

	CustomerID	Highest_Num
1	17841	7983

Part b)

```
# Most Valuable Customer with highest volume of transactions
valuable_customer <- mydf %>% na.omit() %>%
  group_by(CustomerID) %>%
  summarise(Highest_Transaction = sum(TransactionValue)) %>%
  top_n(1, Highest_Transaction)

# Show the dataframe
as.data.frame(valuable_customer)
```

	CustomerID	Highest_Transaction
1	14646	279489

7. Calculate the percentage of missing values for each variable in the dataset. Hint colMeans():

```
# Percentage of missing values
perc_missing_val <- colMeans(is.na(mydf))

# Show the dataframe
as.data.frame(perc_missing_val)
```

	perc_missing_val
InvoiceNo	0.0000000
StockCode	0.0000000
Description	0.0000000
Quantity	0.0000000
InvoiceDate	0.0000000
UnitPrice	0.0000000
CustomerID	0.2492669
Country	0.0000000
TransactionValue	0.0000000
New_Invoice_Date	0.0000000
Invoice_Day_Week	0.0000000
New_Invoice_Hour	0.0000000
New_Invoice_Month	0.0000000

The output shows that Description_name and CustomerID_name have missing values, with 0.27% and 24.93% of missing values respectively.

8. What are the number of transactions with missing CustomerID records by countries?

```
# Total number of transactions with missing CustomerID by country
missing_ID <- mydf %>% group_by(Country, CustomerID) %>%
  filter(is.na(CustomerID)) %>%
  summarise(Num_Transactions = n())

# Show the dataframe
as.data.frame(missing_ID)
```

Country	CustomerID	Num_Transactions
---------	------------	------------------

1	Bahrain	NA	2
2	EIRE	NA	711
3	France	NA	66
4	Hong Kong	NA	288
5	Israel	NA	47
6	Portugal	NA	39
7	Switzerland	NA	125
8	United Kingdom	NA	133600
9	Unspecified	NA	202

```
# To make sure we are counting all the missing CustomerID values
sum(missing_ID$Num_Transactions)
```

```
[1] 135080
```

9. On average, how often the customers comeback to the website for their next shopping? (i.e. what is the average number of days between consecutive shopping) (Optional/Golden question: 18 additional marks!) Hint: 1. A close approximation is also acceptable and you may find `diff()` function useful.

```
# Days average between consecutive shopping
often <- mydf %>% select(CustomerID, New_Invoice_Date) %>%
  group_by(CustomerID) %>%
  mutate(Diff_Days = as.numeric(c(diff(New_Invoice_Date),0))) %>%
  summarise(Time_Days = sum(Diff_Days),
            Days_Avg = sum(Diff_Days)/sum(n()))

# Show the dataframe
head(as.data.frame(often))
```

	CustomerID	Time_Days	Days_Avg
1	12346	0	0.000000
2	12347	365	2.005495
3	12348	283	9.129032
4	12349	0	0.000000
5	12350	0	0.000000
6	12352	260	2.736842

10. In the retail sector, it is very important to understand the return rate of the goods purchased by customers. In this example, we can define this quantity, simply, as the ratio of the number of transactions cancelled (regardless of the transaction value) over the total number of transactions. With this definition, what is the return rate for the French customers?. Consider the cancelled transactions as those where the 'Quantity' variable has a negative value.

```
# Ratio of devolutions from French customers
numenador <- mydf %>% select(Quantity, TransactionValue, Country) %>%
  filter(Country == "France" & Quantity < 0)

denominador <- mydf %>% select(Quantity, TransactionValue, Country) %>%
  filter(Country == "France")

my_ratio <- count(numenador) / count(denominador)

# Show the dataframe
as.data.frame(my_ratio)
```

```
      n
1 0.01741264
```

11. What is the product that has generated the highest revenue for the retailer? (i.e. item with the highest total sum of 'TransactionValue').

```
# Highest revenue for the retailer
highest_revenue <- mydf %>% group_by(Description) %>%
  summarise(highest_transaction = sum(TransactionValue)) %>%
  top_n(1)
```

Selecting by highest_transaction

```
# Show the dataframe
as.data.frame(highest_revenue)
```

```
      Description highest_transaction
1 DOTCOM POSTAGE          206245.5
```

12. How many unique customers are represented in the dataset? You can use unique() and length() functions.

```
# Show the number of unique customers
length(unique(mydf$CustomerID))
```

```
[1] 4373
```